load-diabetes

June 20, 2024

[5]:
```
# Import necessary libraries
from sklearn.datasets import load_diabetes
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

# Load dataset
diabetes = load_diabetes()
X, y = diabetes.data, diabetes.target

# Convert data to DataFrame for easier analysis
df = pd.DataFrame(data=X, columns=diabetes.feature_names)
df['target'] = y

# Display basic statistics and information
print(f"Dataset shape: {df.shape}")
print(f"Columns: {df.columns}")
print(f"Target variable summary:\n{df['target'].describe()}")

# Display correlation heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Heatmap')
plt.show()

# Pairplot for visualizing relationships and distributions
sns.pairplot(df, diag_kind='hist')
plt.suptitle('Pairplot of Diabetes Dataset Features', y=1.02)
plt.tight_layout()
plt.show()

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,␣
↪random_state=42)
```

```python

# Train Linear Regression model
lr = LinearRegression()
lr.fit(X_train, y_train)

# Predict and evaluate
y_pred = lr.predict(X_test)
print(f"\nMean Squared Error: {mean_squared_error(y_test, y_pred)}")
```

```
Dataset shape: (442, 11)
Columns: Index(['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6',
       'target'],
      dtype='object')
Target variable summary:
count    442.000000
mean     152.133484
std       77.093005
min       25.000000
25%       87.000000
50%      140.500000
75%      211.500000
max      346.000000
Name: target, dtype: float64
```

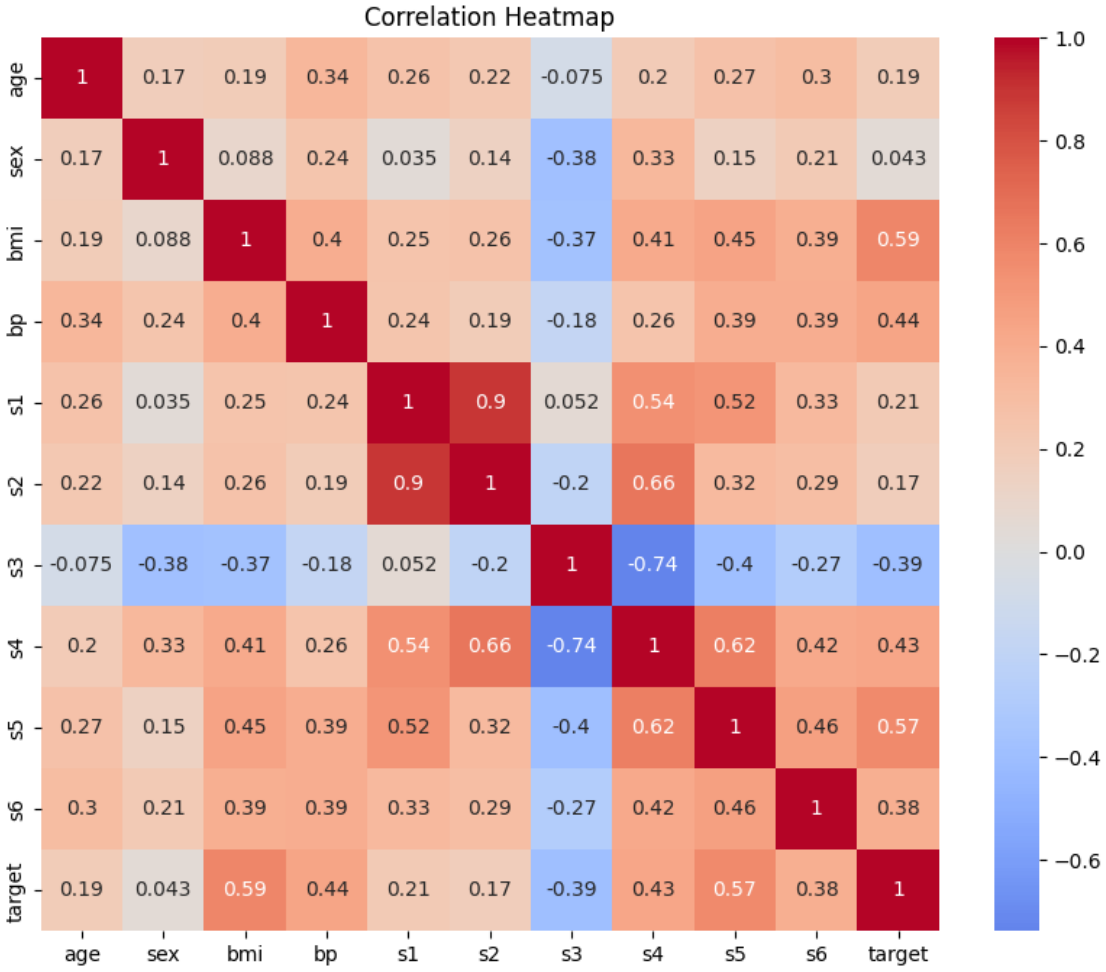
Correlation Heatmap
age 1 0.17 0.19 0.34 0.26 0.22 -0.075 0.2 0.27 0.3 0.19
sex 0.17 1 0.088 0.24 0.035 0.14 -0.38 0.33 0.15 0.21 0.043
bmi 0.19 0.088 1 0.4 0.25 0.26 -0.37 0.41 0.45 0.39 0.59
bp 0.34 0.24 0.4 1 0.24 0.19 -0.18 0.26 0.39 0.39 0.44
s1 0.26 0.035 0.25 0.24 1 0.9 0.052 0.54 0.52 0.33 0.21
s2 0.22 0.14 0.26 0.19 0.9 1 -0.2 0.66 0.32 0.29 0.17
s3 -0.075 -0.38 -0.37 -0.18 0.052 -0.2 1 -0.74 -0.4 -0.27 -0.39
s4 0.2 0.33 0.41 0.26 0.54 0.66 -0.74 1 0.62 0.42 0.43
s5 0.27 0.15 0.45 0.39 0.52 0.32 -0.4 0.62 1 0.46 0.57
s6 0.3 0.21 0.39 0.39 0.33 0.29 -0.27 0.42 0.46 1 0.38
target 0.19 0.043 0.59 0.44 0.21 0.17 -0.39 0.43 0.57 0.38 1
age sex bmi bp s1 s2 s3 s4 s5 s6 target
1.0 0.8 0.6 0.4 0.2 0.0 −0.2 −0.4 −0.6

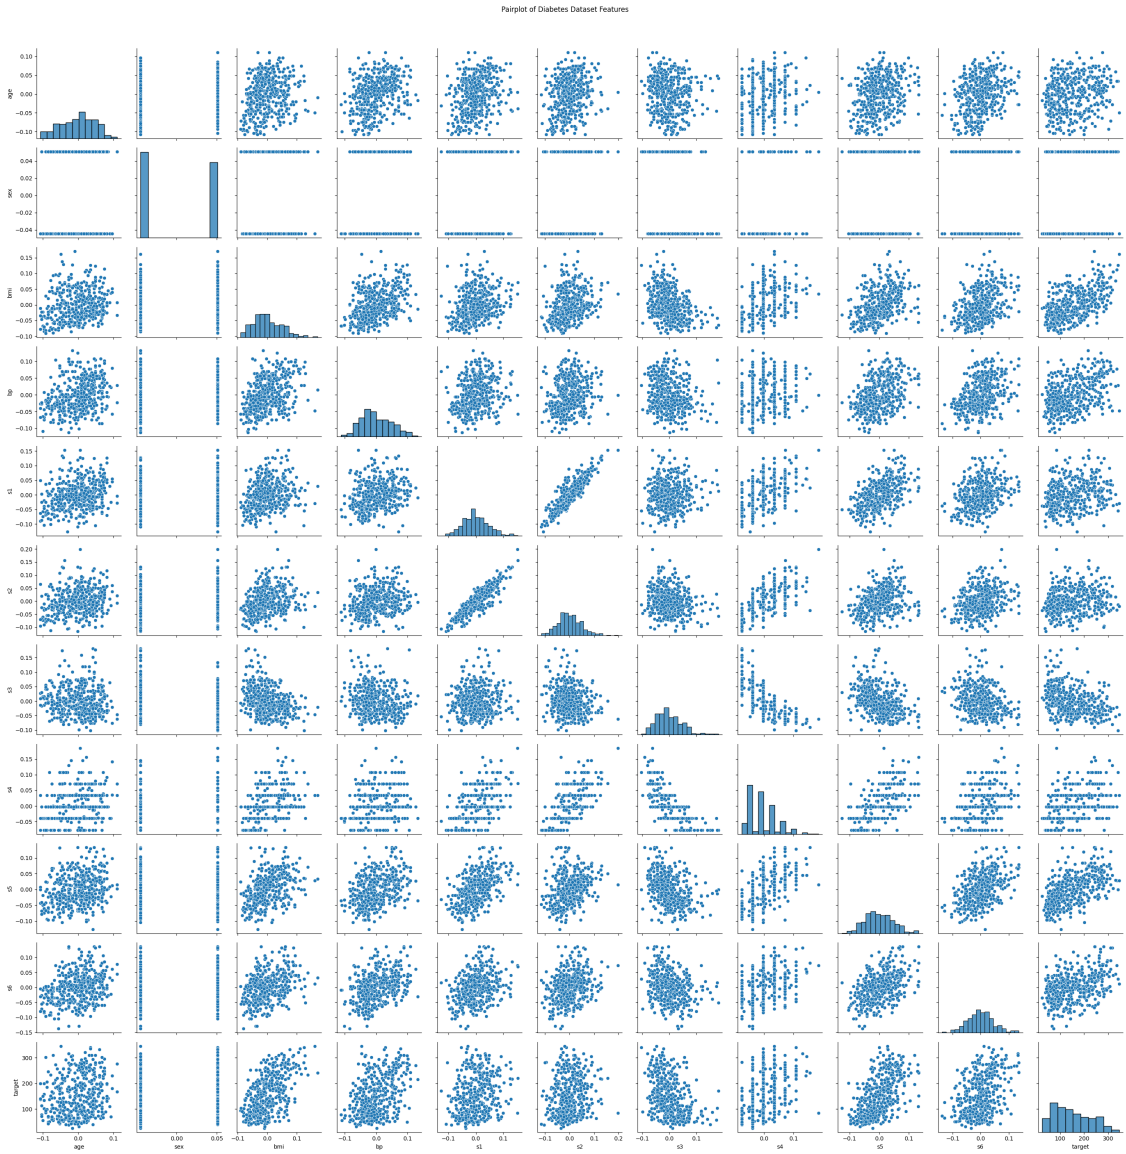


```
Mean Squared Error: 2900.193628493482
```